

1. What We Have Now

- Our current system can detect brain tumors from MRI scans using AI (ResNet18 model).
 - Users upload images on a React web app; the Flask backend predicts the tumor type, gives a basic explanation (caption), and shows a Grad-CAM heatmap for visual proof.
-

2. What We're Adding Next

A. Fracture Detection

- We will train a YOLO model to detect bone fractures in X-ray images.
- The system will work similarly: user uploads image, backend detects fracture, and displays results.

B. Smarter Captions with GenAI

- Instead of fixed captions, we'll use advanced AI (LLMs) to write clear, personalized medical explanations for both tumor and fracture results.

C. Severity Assessment

- The AI will also rate how serious the tumor or fracture is, and give simple follow-up advice.

D. RAG-based Report Summaries & Chat

- Users can upload medical reports (PDFs); our system will summarize them and answer questions by finding info in those PDFs.

E. Chatbot with Image Memory

- Users can chat with an AI assistant about their images and reports.
- The bot remembers previous uploads and can compare old and new results.

F. Proactive Clinical Assistant (Agentic AI)

- The assistant will not just answer but also suggest next steps:
 - For example, if a high-risk tumor is found, it will offer to show clinical guidelines or help schedule a doctor's visit.
 - If no fracture is found, but the user has pain, it may suggest uploading a follow-up image later.

1. Fracture Detection via YOLO

- **Technical Detail:**

- We will train a YOLO model (You Only Look Once) on a fracture X-ray dataset.
 - The Flask backend will load this new model file (e.g. yolo_fracture.pth) and add a new API endpoint (like /fracture/predict).
 - The frontend will let the user upload an X-ray, send it to the backend, and display bounding boxes showing detected fractures.
-

2. LLM-Based Captioning for MRI & Fracture

- **Technical Detail:**

- Instead of showing fixed captions, we'll use a Large Language Model (LLM) like GPT or Llama.
 - This model will take results (tumor/fracture type, image region) and generate a custom, easy-to-understand medical report for each scan.
 - The backend will add an endpoint (e.g. /caption/generate) that sends data to the LLM and returns the generated text.
-

3. Severity Assessment with LLM

- **Technical Detail:**

- After detection, the AI will rate how serious the tumor or fracture is.
 - The backend will use LLM prompts based on model results to give a severity score (“low”, “moderate”, “high”) and explain why.
 - The frontend will show this score alongside the results.
-

4. RAG-Based Clinical Report Summaries & Chat

- **Technical Detail:**

- Users can upload PDFs of medical reports.
- The backend will use Retrieval-Augmented Generation (RAG): it finds relevant info from the report using embeddings and returns LLM-generated summaries or answers.

- A new API (e.g. /rag/query) will handle these requests, and the frontend will show answers with document citations.
-

5. Chatbot with Image Memory

- **Technical Detail:**

- Users chat with an AI assistant about their scans.
 - The backend keeps session memory: it can recall previous images and results, so users can ask things like “compare this to my last scan.”
 - The frontend will have a chat interface that connects to an endpoint like /chat/message.
-

6. Proactive Agent Clinical Assistant

- **Technical Detail:**

- The AI agent monitors results and user profile.
 - If a high-risk result is found, it can automatically suggest next steps (like “view guidelines,” “schedule appointment,” or “upload old scans for comparison”).
 - This logic runs in the backend and triggers prompts in the frontend for user action—always requiring user consent before taking any step.
-

7. Multi-Model Backend Structure

- **Technical Detail:**

- Each model (tumor, fracture, captions) will have its own Python file and Flask endpoint.
- For example, model_mri.pth for tumor, yolo_fracture.pth for fracture.
- The backend will load these models separately and route user requests to the right endpoint, making it easy to add new analysis types in the future.

1. Frontend (React)

User Experience

- **Login & Profile:**
 - User signs in (see Login.tsx), profile info managed (see AuthContext.tsx and Profile.tsx).
 - Can edit details, upload profile picture, and store medical history for personalized analysis.
 - **Image Upload & Analysis:**
 - User uploads an MRI or X-ray on ImageUpload.tsx.
 - The page shows image preview, processing animation, and then results (prediction, explanation, Grad-CAM image).
 - **Result Display:**
 - Displays predicted label (tumor/fracture type), AI-generated explanation, and visual Grad-CAM (for MRI) or bounding boxes (for fractures).
 - **Interactive Chat & RAG Features:**
 - User can chat about their scan or upload medical PDFs for instant summaries and Q&A.
-

2. Backend (Flask)

Core API Structure

- **Modular Endpoints:**
 - /predict (MRI): Uses ResNet18, returns tumor label, caption, Grad-CAM.
 - /fracture/predict: Will use YOLO, return fracture type and bounding boxes.
 - /caption/generate: Uses LLM to make a readable report/caption based on findings.
 - /severity/score: LLM rates seriousness (low/high) and gives advice.
 - /rag/query: For document-based answers using clinical PDFs.
 - /chat/message: Smart assistant, keeps chat session and responds based on current and previous scans.

- `/agent/suggest`: Agentic AI logic, proactively suggests actions based on findings.

Model Handling

- **Multiple model files:**

- Each task (MRI, fracture) has its own .pth model loaded once at server start.
 - Model registry or config file points Flask to correct weights for each endpoint.
 - Easy to add new models: just update registry/config and add a new endpoint.
-

3. Medical AI Models

A. MRI Tumor Detection

- **Model:** ResNet18, fine-tuned for 4 classes (glioma, meningioma, pituitary, notumor).
- **Explainability:** Grad-CAM highlights suspicious regions, shown to user.

B. Fracture Detection

- **Model:** YOLOv8 (or latest), trained on bone fracture X-rays.
- **Output:** Bounding boxes on the image, fracture type, confidence score.

C. Future Models

- **KAN, other CNNs:** For experiments, or to improve accuracy.
-

4. GenAI Modules

A. Caption Generation (LLM)

- **LLM Model:** GPT-4o, Llama2, or similar (hosted via API or local inference).
- **Input:** Structured findings (tumor/fracture type, region, severity, patient history).
- **Output:**
 - Patient-friendly text: “Your MRI shows a small glioma in the left temporal lobe. No signs of spread.”
 - Clinician summary: “Glioma, left temporal, 2.3cm, no edema.”

- **Integration:** Flask endpoint /caption/generate calls the LLM, returns result to frontend.

B. Severity Assessment

- **LLM uses rules:**
 - Reads size, location, displacement, confidence, patient history.
 - Outputs severity (e.g., “High risk due to location near motor cortex”) and advice.
- **Endpoint:** /severity/score

C. RAG (Retrieval-Augmented Generation)

- **PDF Upload:** User uploads clinical reports/guidelines.
- **How it works:**
 - Backend splits PDF into chunks, creates embeddings, stores in FAISS (vector DB).
 - When user asks a question, backend finds relevant chunks and LLM answers only using info from those chunks.
 - Citations: Each answer links to the exact PDF page/section.
- **Endpoints:** /rag/ingest, /rag/query

D. Chatbot with Memory

- **LLM-powered conversation:**
 - Backend stores session data (previous images, findings, chat history).
 - When user chats, LLM gets current findings and session history for smarter replies.
- **Example:** “Last scan showed a meningioma. This one looks normal. Want to compare images?”
- **Endpoint:** /chat/message

E. Agentic Clinical Assistant

- **Logic:**
 - Monitors predictions and user profile.
 - Proactively suggests actions (not just answers questions).
- **Examples:**

- “High-grade glioma detected. Would you like to view treatment guidelines or schedule a specialist visit?”
 - “No fracture found, but your profile says ongoing pain. Should we set a reminder to check again in two weeks?”
 - **Endpoint:** /agent/suggest
 - **Opt-in:** User must accept before actions (privacy-first).
-

5. Backend Tech Stack

- **Flask:** Serves all AI endpoints, handles API requests from frontend.
 - **PyTorch:** Loads and runs ResNet/YOLO models for predictions.
 - **FAISS:** Fast vector search for RAG.
 - **Session store:** SQLite or Redis for chat and user history.
 - **Joblib:** Loads label encoders and other pickled utilities.
-

6. Frontend Tech Stack

- **ReactJS:** SPA for all pages.
 - **Tailwind CSS:** For beautiful, responsive design.
 - **React Dropzone:** For image upload UI.
 - **LocalStorage:** For storing user/session data.
 - **REST API calls:** To Flask endpoints for predictions, captions, chat, etc.
-

7. How All Features Work Together (Flow)

1. **User logs in, completes profile.**
2. **Uploads an MRI or fracture X-ray.**
3. **Frontend sends image to backend.**
4. **Backend runs the correct model (ResNet18 or YOLO), gets prediction.**
5. **Grad-CAM/boxes created for visual explanation.**
6. **Prediction sent to LLM for easy-to-read caption/report.**
7. **LLM also scores severity and gives advice.**

8. **User can chat about findings, compare with old scans, or upload PDFs for instant Q&A.**
 9. **Agentic assistant monitors, and if something important is found, proactively suggests next steps (with user consent).**
 10. **All results, chats, and history stored for future reference.**
-

8. Adding New Models (Technical Plan)

- Put each new model (e.g., model_fracture.pth) in a models/ folder.
 - Register it in a config file, mapping task → model file.
 - Add a new Flask endpoint (e.g., /fracture/predict) that loads and uses only that model.
 - Frontend adds a new upload mode (MRI or fracture), and calls the right endpoint based on image type.
 - Results handled in the same way—modular, easy to scale.
-

9. Deployment & Scaling

- **Backend:** Runs on cloud VM or college server, with CUDA for speed if available.
 - **Frontend:** Hosted on Vercel/Netlify or college web server.
 - **Security:** CORS enabled, authentication required, user data protected.
 - **Models:** Loaded once at startup, fast inference (CPU or GPU).
 - **RAG:** PDF data stored locally; embeddings indexed for instant Q&A.
-

10. Why This Stack?

- **Modular:** Easy to add new models and features.
 - **Scalable:** Can handle more tasks (other diseases, more imaging types).
 - **User-friendly:** Clear reports, chat, and proactive help.
 - **AI-powered:** Combines classic deep learning (for detection) with cutting-edge GenAI (for captions, chat, RAG, agentic logic).
-

11. Summary Table (Feature → Tech Details)

Feature	Tech Details
MRI Tumor Detection	ResNet18, Grad-CAM, Flask /predict endpoint
Fracture Detection	YOLOv8, Flask /fracture/predict endpoint
Caption Generation	LLM (GPT-4o/Llama2), Flask /caption/generate
Severity Assessment	LLM + prompt, Flask /severity/score
RAG-based Q&A	PDF chunking, SBERT/MiniLM embeddings, FAISS, LLM, Flask /rag/query
Chatbot	LLM, session/context store, Flask /chat/message
Agentic Assistant	Trigger logic + LLM, Flask /agent/suggest
Frontend	ReactJS, Tailwind, REST calls, pages: upload, results, chat, profile
Adding Models	Place new .pth file, update config, add endpoint, update frontend mode selector